

San Francisco Real-Time Traffic Prediction & Optimization Project

Project Mission

To build a cloud-based system that uses real traffic data from San Francisco to predict congestion and dynamically regulate traffic light timings at intersections, reducing delays and improving urban mobility through AI-driven decision-making.

Phase 1: Data Acquisition & Road Network Setup

Goal: Prepare a clean, queryable dataset and road graph for San Francisco using real-world data.

Steps:

1. GCP Environment Setup

- Create GCP project
- Enable APIs: BigQuery, Vertex AI, Cloud Run, Maps API, Cloud Storage
- Create BigQuery dataset (e.g., `sf_traffic_data`)
- Create service account with appropriate roles

2. Download & Inspect Data

- Traffic Flow Segments (SF Open Data)
- Signalized Intersections (SF Open Data)
- Uber Movement SF (optional, for historical patterns)

3. Upload to BigQuery

- `traffic_segments`
- `signal_intersections`
- `uber_travel_times` (optional)

4. Build Road Network Graph

- Use NetworkX in Python
- Nodes = intersections

- Edges = road segments with attributes (speed, geometry, time windows)

5. Notebook & Exploration

- Create Jupyter notebook for data exploration and graph creation

Deliverables: BigQuery tables, Python road graph object, EDA notebook

Phase 2: Train Congestion Prediction Model

Goal: Predict traffic congestion for segments using real historical data.

Steps:

1. Preprocessing

- Aggregate speeds by segment and time window
- Generate features: time, speed, length, previous readings

2. Model Selection

- Choose between XGBoost, Random Forest (tabular) or LSTM (time series)

3. Training & Evaluation

- Train and validate with train-test split (avoid time leakage)
- Measure MAE, RMSE, etc.

4. Model Export

- Save with `joblib` or `SavedModel`
- Upload to GCS

Deliverables: Trained model + evaluation report

Phase 3: Define Traffic Light Timing Logic

Goal: Adjust signal timings based on predicted congestion.

Steps:

1. Baseline Setup

- Fixed-cycle signal timings for control

2. Adaptive Logic Design

- Extend green/red based on predicted congestion
- Optionally implement queue-length or throughput-based rules

3. RL Option (Advanced)

- Reinforcement learning agent (e.g., DQN) to control signals

4. Simulation

- Simulate using prediction + logic loop
- Track metrics: delay, queue length, throughput

Deliverables: Signal logic engine + performance metrics

Phase 4: Build & Deploy Backend (Cloud Run + Vertex AI)

Goal: Host prediction model and signal logic as APIs.

Steps:

1. Model Deployment

- Deploy to Vertex AI or serve via Flask on Cloud Run

2. API Development

- Endpoints: `/predict_congestion`, `/signal_state`, `/update_traffic_state`

3. Cloud Setup

- Configure service account, logging, billing
- Add Stackdriver monitoring

Deliverables: Deployed backend with public REST API

Phase 5: Build and Connect the Dashboard

Goal: Create a web interface to visualize live and predicted traffic data.

Steps:

1. Choose Framework

- Streamlit (fast prototyping) or React + Mapbox

2. Implement Features

- Map of SF
- Predicted congestion (color-coded)
- Signal light states at intersections

3. Backend Integration

- Fetch predictions and logic outputs via REST API

4. Live Controls (Optional)

- Allow users to simulate scenarios or switch modes

Deliverables: Web dashboard connected to backend

Phase 6: Documentation & Project Packaging

Goal: Make the project clear, reproducible, and shareable.

Steps:

1. Technical Docs

- Describe architecture, data flow, ML model, cloud setup
- Use diagrams, flowcharts, and inline comments

2. Code Docs & README

- Add `README.md` with setup, usage, structure, and credits
- Add docstrings and inline comments to code

3. Experiment Logs

- Save performance comparisons of signal strategies
- Include visualizations: delay charts, speed heatmaps

4. Reproducibility Setup

- Add `requirements.txt`, `setup.sh`, Dockerfile, or notebooks
- Add quickstart and teardown scripts for GCP

5. Final Assets

- Upload UI screenshots, architecture diagrams, optional demo video
- Include LICENSE and acknowledgment section

Deliverables: Full project documentation + packaging for deployment or handoff